

Nirbhay sharma (B19CSE114)

Regularizing Neural Networks via Adversarial Model Perturbations

How to Run

Installing requirements

```
pip install -r requirements.txt
```

Run for CIFAR10 dataset

```
CUDA_VISIBLE_DEVICES=<gpu_id> python main.py --dataset cifar10 --model  
<model_name>
```

Run for gtsrb dataset

```
CUDA_VISIBLE_DEVICES=<gpu_id> python main.py --dataset gtsrb --model  
<model_name>
```

Execution flow of the code

The algorithm is outlined below.

Algorithm 1 Adversarial Model Perturbation Training

Require: Training set $\mathcal{D} = \{(\mathbf{x}, \mathbf{y})\}$, Batch size m , Loss function ℓ , Initial model parameter θ_0 , Outer learning rate η , Inner learning rate ζ , Inner iteration number N , L_2 norm ball radius ϵ

```

1: while  $\theta_k$  not converged do
2:   Update iteration:  $k \leftarrow k + 1$ 
3:   Sample  $\mathcal{B} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^m$  from training set  $\mathcal{D}$ 
4:   Initialize perturbation:  $\Delta_{\mathcal{B}} \leftarrow \mathbf{0}$ 
5:   for  $n \leftarrow 1$  to  $N$  do
6:     Compute gradient:
        $\nabla \mathcal{J}_{\text{AMP}, \mathcal{B}} \leftarrow \sum_{i=1}^m \nabla_{\theta} \ell(\mathbf{x}_i, \mathbf{y}_i; \theta_k + \Delta_{\mathcal{B}}) / m$ 
7:     Update perturbation:  $\Delta_{\mathcal{B}} \leftarrow \Delta_{\mathcal{B}} + \zeta \nabla \mathcal{J}_{\text{AMP}, \mathcal{B}}$ 
8:     if  $\|\Delta_{\mathcal{B}}\|_2 > \epsilon$  then
9:       Normalize perturbation:  $\Delta_{\mathcal{B}} \leftarrow \epsilon \Delta_{\mathcal{B}} / \|\Delta_{\mathcal{B}}\|_2$ 
10:    end if
11:  end for
12:  Compute gradient:
     $\nabla \mathcal{J}_{\text{AMP}, \mathcal{B}} \leftarrow \sum_{i=1}^m \nabla_{\theta} \ell(\mathbf{x}_i, \mathbf{y}_i; \theta_k + \Delta_{\mathcal{B}}) / m$ 
13:  Update parameter:  $\theta_{k+1} \leftarrow \theta_k - \eta \nabla \mathcal{J}_{\text{AMP}, \mathcal{B}}$ 
14: end while

```

The whole problem can be seen as a minimization problem of the $\max_{\delta \in B(\mu, \epsilon)} (L_{EMP}(x, y, \theta + \delta))$ so the first task is find a suitable δ which lies in L_2 ball with radius ϵ and then simply minimizing it over model parameters θ would do the work.

so the entire algorithm is as follows:

- first initialize the parameters (θ, δ)
- run for specific iterations (N) to obtain δ using gradient ascent
 - calculate gradient of loss as $\sum \nabla_{\theta} l(x, y; \theta + \delta)$
 - update δ as $\delta = \delta + \eta \sum \nabla_{\theta} l(x, y; \theta + \delta)$
- once δ is obtained again compute the gradient using the same formula and update the model parameters as follows
 - $\theta = \theta - \eta \sum \nabla_{\theta} l(x, y; \theta + \delta)$
- The entire workflow first finds a suitable delta in L_2 ball and then update the model parameters accordingly.

The authors have implemented this as an optimization algorithm known as AMP in the code as follows.

```
@torch.no_grad()
def step(self, closure=None):
    if closure is None:
        raise ValueError('Adversarial model perturbation requires closure, but it was not provided')
    closure = torch.enable_grad()(closure)
    outputs, loss = map(lambda x: x.detach(), closure())
    for i in range(self.defaults['inner_iter']):
        for group in self.param_groups:
            for p in group['params']:
                if p.grad is not None:
                    if i == 0:
                        self.state[p]['dev'] = torch.zeros_like(p.grad)
                        dev = self.state[p]['dev'] + group['inner_lr'] * p.grad
                        clip_coef = group['epsilon'] / (dev.norm() + 1e-12)
                        dev = clip_coef * dev if clip_coef < 1 else dev
                        p.sub_(self.state[p]['dev']).add_(dev) # update "theta" with "theta+delta"
                        self.state[p]['dev'] = dev
                    closure()
        for group in self.param_groups:
            for p in group['params']:
                if p.grad is not None:
                    p.sub_(self.state[p]['dev']) # restore "theta" from "theta+delta"
    self.base_optimizer.step()
    return outputs, loss
```

Here authors have done the same thing as described above, first they find out the suitable delta to use as a regularizer and then they use standard optimization algorithm (SGD) to finally update the model parameters.